Subhajit Chowdhury

Cloud Data Engineer

Kolkata, India | +91 7029860941 | er.subhajitchowdhury@gmail.com

linkedin.com/in/subhajit-dataengineer | github.com/Subhajit-Chowdhury | subhajit-chowdhury.github.io

Professional Summary

Cloud Data Engineer with nearly 4 years of experience specializing in AWS architectures, distributed processing, and enterprise cloud modernization. Architected and maintained 150+ production-grade ETL/ELT pipelines processing hundreds of GBs daily with periodic TB-scale refreshes across finance and retail domains. Proficient in PySpark, Snowflake, and Medallion architectures (Delta Lake), with a proven track record of reducing analytics data latency by 30% and generating £487K in annual infrastructure savings through strategic AWS workload optimization.

Technical Skills

Cloud & Data Platforms: AWS (S3, Glue, Lambda, Athena, Redshift, Step Functions, DynamoDB, EMR), Snowflake
Data Engineering: ETL/ELT, Data Lakes, Medallion Architecture, Delta Lake, CDC, Informatica IICS
Programming & Querying: Python (Pandas, NumPy), PySpark, SQL (PostgreSQL, Oracle, Redshift), Shell Scripting
Orchestration & DevOps: AWS Step Functions, Jenkins, Git, CloudWatch, CI/CD, Agile/Scrum

Professional Experience

Tata Consultancy Services (TCS) — July 2022 – Present
Data Engineer

Client: Aegon (Global) — *July 2024 – Present*

- Architected and maintained **150+ production-grade AWS ETL/ELT pipelines** (S3, Glue, EMR, Redshift, Athena) across a medallion architecture (Bronze/Silver/Gold), processing hundreds of GBs daily with periodic **TB-scale refreshes** for downstream BI consumers
- Standardized an ACID-compliant **Medallion** data lake on **Delta Lake** and **Parquet**
- Automated schema discovery using **AWS Glue Crawlers**, cutting downstream data latency by **30%**
- Optimized **Redshift and Athena** SQL queries by **20–25%** via partition pruning and predicate pushdowns
- Developed incremental load logic (insert/update/delete) using **PySpark** and **SQL** for dimensional consistency
- Established CI/CD pipelines using **Jenkins and Git** for version-controlled, repeatable production deployments

Client: Kingfisher plc (Europe) — *July 2023 – June 2024*

- Migrated on-premises workloads to AWS, deploying modernized cloud features to save **£487K** annually
- Orchestrated fault-tolerant batch ETL workflows via **AWS Step Functions, Glue, and Lambda**
- Engineered scalable delivery pipelines routing curated datasets to **6 distinct downstream BI systems**
- Utilized **AWS DMS** for continuous automated data ingestion into Amazon S3 raw storage layers

Client: McGraw Hill (Global) — *July 2022 – June 2023*

- Executed on-premises Oracle database migration to **Oracle Cloud (OCI)** ensuring zero data loss
- Migrated complex data integration layers from legacy Informatica PowerCenter to **Informatica IICS**
- Developed end-to-end ETL mappings utilizing **Oracle SQL** for robust data cleansing and extraction

Technical Projects

Near-Real-Time Serverless ELT Pipeline | *DynamoDB, Lambda, Snowflake* | GitHub

- Engineered an event-driven pipeline capturing live API data using **DynamoDB Streams and Lambda**
- Automated continuous S3 data ingestion into Snowflake via **Snowpipe**, bypassing batch schedules

Serverless Batch Analytics ETL Pipeline | *Amazon S3, AWS Glue, PySpark, Athena* | GitHub

- Built a data lake incorporating schema discovery and distributed **PySpark** transformations via AWS Glue
- Integrated AWS Glue Data Catalog with **Amazon Athena** for direct SQL querying on S3-stored data

Education

Bachelor of Technology(B.Tech) in Computer Science and Engineering — 2018 – 2022
Modern Institute of Engineering and Technology (MAKAUT) — *GPA: 8.90*

Certifications

AWS Certified Cloud Practitioner | a6447414ae25419cba12941481cd10a2 | *Validation: July 2025 – July 2028*